

Introduction to OpenMP

A guide by WITS HPC

Unleashing parallel power

We are limited by single processor chips:

- As data sizes and processing demands grow, single processors struggle to keep up.
- Sequential programs execute instructions one after another, leading to bottlenecks





Enter OpenMP

What is parallel computing?

→ Breakdown

Parallel programming breaks down large tasks into smaller, independent subtasks.

→ Execution

These subtasks can be executed simultaneously on multiple processors (cores) or computers.





OpenMP: The Message Passing Maestro

What even is OpenMPI?

→ What

Open Message Passing Interface (OpenMPI) is a free, open-source library for parallel programming.

→ How

It provides a standardized set of functions for processes to communicate and exchange data.

So how does it work?

→ Create

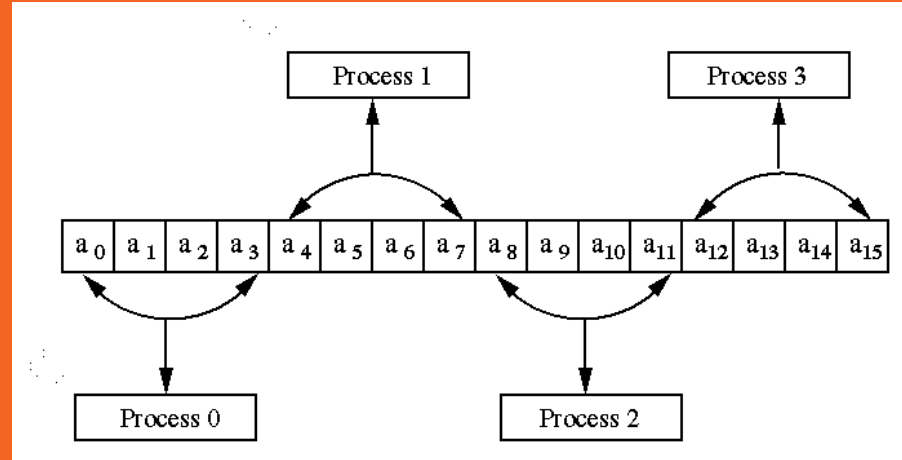
OpenMPI creates multiple processes (copies of your program).

→ Communicate

Processes reside on different processors or computers and communicate by sending and receiving messages.

→ Partitioned

Data is partitioned and distributed among processes.



#pragma omp parallel

- Creates a team of threads for parallel execution of the following code block.

#pragma omp for

- Distributes iterations of a loop across threads for parallel execution.

#pragma omp sections

- Defines multiple independent sections of code to be executed by different threads.

#pragma omp single

- Specifies a code block to be executed by only one thread.

#pragma omp master

- Specifies a code block to be

MP Directives.

All MPI sections start with

#pragma omp <construct>
<clauses>

{

//code block to run in parallel

}

Some MP clauses.

These give us more control over the program.

Story for illustration purposes only

private(list)

- Creates a separate copy of listed variables for each thread.

shared(list)

- Makes listed variables accessible by all threads.

firstprivate(list)

- Like private, but initializes each copy with the value of the original variable.

reduction(operator: list)

- Performs a reduction operation on variables at the end of the parallel region (e.g., +, *, max, min).

nowait

- Removes the implicit barrier at the end of a parallel construct, allowing threads to continue without waiting.
-

Basic hello world program

```
#include <iostream>
#include <omp.h>

using namespace std;

int main(){
    int num_threads = omp_get_num_threads(); //used to get the number of threads available
    cout << "We have a total of: " << num_threads << " threads available!" << endl;
    cout << "We only have one since we are running outside of a omp parallel region" << endl;
    #pragma omp parallel //define an omp parallel region
    {

        int ID = omp_get_thread_num();
        cout << "Hello from thread: " << ID << endl; //race condition!

    }

    return 0;
}
```

To compile :

g++ -fopenmp <filename.cpp> -o execname


```

#include <iostream>

using namespace std;

static long num_steps = 1000000;
double step;

int main(){
    int i;
    double x, pi, sum = 0.0;

    step = 1.0/( double ) num_steps ;

    for(i=0; i < num_steps; ++i){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }

    pi = step * sum;
    cout << "sequential: pi with " << num_steps << " is: " << pi << endl;
    return 0;
}

```

Concrete example of why MPI is better

Here is a simple program to find the sequential sum

The serial runs in:

```
anandpatel@pop-os:~/University/PC/week-7$ time ./pi_serial_intel
sequential: pi with 1000000 is: 3.14159
```

```
real    0m0.011s
user    0m0.007s
sys     0m0.004s
```

Whereas in parallel we get:

```
anandpatel@pop-os:~/University/PC/week-7$ ./pi_par_intel
running on 1 threads: PI = 3.141592653589903 computed in 0.003019 seconds
running on 2 threads: PI = 3.141592653589862 computed in 0.0017 seconds
running on 3 threads: PI = 3.141592653589884 computed in 0.0011 seconds
running on 4 threads: PI = 3.141592653589873 computed in 0.0008969 seconds
running on 5 threads: PI = 3.141592653589874 computed in 0.001328 seconds
running on 6 threads: PI = 3.141592653589876 computed in 0.001716 seconds
running on 7 threads: PI = 3.141592653589876 computed in 0.001669 seconds
running on 8 threads: PI = 3.141592653589878 computed in 0.001725 seconds
```

